Blind judgement (by a non-contestant model)

Response	Criterion	Score	One-sentence justification
A	Clarity	8	Headings, lists, and the worked broadcast example make the argument accessible, although the proliferation of subsections fragments the document.
A	Concision	5	The LDFI loop, certificate scope, coverage limits, and dissemination argument recur in multiple sections without adding much each time.
A	Precision	7	Concrete fault descriptions and examples help, but “pump in faults until a labelled collection cracks” and calling sampled results “Certified” weaken the terminology.
A	Fitness for purpose	7	It explains complete EC and classifies existing code explicitly, but its playful subtitle and extensive expansion work against the requested direct engineering style.
A	Correctness	7	Most implementation details match the source, but grouping sampled searches under certification contradicts its own caveats, and the worked echo example asserts an unverified search sequence.
B	Clarity	7	The backwards-search explanation is readable, but several dense paragraphs retain the original draft’s nested qualifications and compressed technical phrasing.
B	Concision	6	It is more compact than A, yet phrases such as “The division of labour is worth stating plainly” and repeated search explanations add unnecessary framing.
B	Precision	6	Specific APIs and fault examples are useful, but claims that certificates cover every element and that sampled protocols are certified exceed the stated check.
B	Fitness for purpose	7	It addresses motivation, operation, and classification, but the decorative opening and retained formulations such as “no form argument grades them” do not sufficiently escape the original style.

Response	Criterion	Score	One-sentence justification
B	Correctness	6	It reproduces the source’s conflict between sampled results and certificates without resolving it, leaving the central classification misleading.
C	Clarity	9	It presents the missing EC argument early, explains the check through a simple contents table, and consistently separates fault search from schedule exploration.
C	Concision	7	Some scope and limitation statements repeat, but most of the added structure clarifies distinctions rather than merely restating them.
C	Precision	9	It distinguishes bounded convergence, probe-domain agreement, exhaustive certification, sampled evidence, and unsupported claims using concrete, consistent terms.
C	Fitness for purpose	9	Its restrained technical register and explicit explanation of complete EC closely match the needs of a Hydro-familiar engineering audience.
C	Correctness	9	It resolves the draft’s misleading certification labels and qualifies the monotonicity argument instead of treating a fixed fact domain as sufficient by itself.

Ranking

1. **C**
2. **A**
3. **B**

C wins because it makes the document’s central distinctions clear and keeps them consistent: conditional type-level EC versus fault-tolerant dissemination, bounded failure versus unbounded nonconvergence, and certificates versus successful samples. A is second because its lists and worked example improve first-pass comprehension, though it repeats too much and undermines its own caveats with the “Certified” grouping. B offers a good intuitive LDFI explanation, but retains more of the original’s compressed, ambiguous prose and leaves the certification conflict unresolved.